

Multi Agent Systems

- Lab 6 -

N-Step Bootstrapping

Recap: state-value prediction

In estimating the value of a policy v_π there are two extremes:

- The **return** G_t of a state reward sequence is
 - $G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T$
 - G_t is the *target* in MC updates
 - $V^\pi(s) = E_\pi[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s] = E_\pi[G_t | S_t = s]$
- In *one-step* updates (Value-Iteration; Q-Learning, SARSA – based on TD learning) the *target* is first reward + discounted estimate for value of next state
 - $G_{t:t+1} \doteq R_{t+1} + \gamma V_t(S_{t+1})$

State-value prediction generalization

- The ***n*-step return** $G_{t:t+n}$ of a state reward sequence is

- $G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n})$
- $V_{t+n}(S_t) \doteq V_{t+n-1}(S_t) + \alpha [G_{t:t+n} - V_{t+n-1}(S_t)], \quad 0 \leq t < T$

n-step TD for estimating $V \approx v_\pi$

Initialize $V(s)$ arbitrarily, $s \in \mathcal{S}$

Parameters: step size $\alpha \in (0, 1]$, a positive integer n

All store and access operations (for S_t and R_t) can take their index mod n

Repeat (for each episode):

Initialize and store $S_0 \neq \text{terminal}$

$T \leftarrow \infty$

For $t = 0, 1, 2, \dots$:

| If $t < T$, then:

| Take an action according to $\pi(\cdot | S_t)$

| Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

| If S_{t+1} is terminal, then $T \leftarrow t + 1$

| $\tau \leftarrow t - n + 1$ (τ is the time whose state's estimate is being updated)

| If $\tau \geq 0$:

| $G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

| If $\tau + n < T$, then: $G \leftarrow G + \gamma^n V(S_{\tau+n})$ ($G_{\tau:\tau+n}$)

| $V(S_\tau) \leftarrow V(S_\tau) + \alpha [G - V(S_\tau)]$

Until $\tau = T - 1$

N-step SARSA

- Use of N-step methods for control as well, besides prediction
- N-step method + SARSA → on-policy TD control method
 - $G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q_{t+n-1}(S_{t+n}, A_{t+n}), \quad n \geq 1, 0 \leq t < T - n$
 - $Q_{t+n}(S_t, A_t) \doteq Q_{t+n-1}(S_t, A_t) + \alpha [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)], \quad 0 \leq t < T$

N-step SARSA

n-step Sarsa for estimating $Q \approx q_*$, or $Q \approx q_\pi$ for a given π

Initialize $Q(s, a)$ arbitrarily, for all $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize π to be ε -greedy with respect to Q , or to a fixed given policy

Parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$, a positive integer n

All store and access operations (for S_t , A_t , and R_t) can take their index mod n

Repeat (for each episode):

 Initialize and store $S_0 \neq$ terminal

 Select and store an action $A_0 \sim \pi(\cdot | S_0)$

$T \leftarrow \infty$

 For $t = 0, 1, 2, \dots$:

 If $t < T$, then:

 Take action A_t

 Observe and store the next reward as R_{t+1} and the next state as S_{t+1}

 If S_{t+1} is terminal, then:

$T \leftarrow t + 1$

 else:

 Select and store an action $A_{t+1} \sim \pi(\cdot | S_{t+1})$

$\tau \leftarrow t - n + 1$ (τ is the time whose estimate is being updated)

 If $\tau \geq 0$:

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

 If $\tau + n < T$, then $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$ ($G_{\tau:\tau+n}$)


$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha [G - Q(S_\tau, A_\tau)]$

 If π is being learned, then ensure that $\pi(\cdot | S_\tau)$ is ε -greedy wrt Q

 Until $\tau = T - 1$

Test environment and Task

- FrozenLake-4x4 environment in OpenAI Gym
- **Task:** analyze the state-value prediction accuracy of TD(0) methods (Q-Learning, SARSA) and TD(n) methods (n-step SARSA)



Task steps

- Implement **n-step SARSA** agent for the **Frozen Lake - small (4x4)** environments
- **Run Value Iteration** on the environment to compute the *ground truth* state-value function – obtain $V^*(s)$
- For **alpha** = [0.0, 0.1, 0.3, 0.5, 0.7, 1.0]
 - For **20 repetitions** with the chosen **alpha**
 - Initialize Q values to 0
 - Run **Q-Learning** and **SARSA** for **8000** episodes
 - Run **n-step SARSA** for 8000 episodes, where $n = 2, 3$ and 4
 - At the end of each episode compute the RMSE (root mean squared error) between $V^*(s)$ and $\max_a Q(s, a)$
 - After learning for the 8000 episodes in each repetition, average MSE over the **20 repetitions**
- **On a same graph**, plot RMSE errors for Q-Learning, SARSA and n-step SARSA (for each value of n): **x-axis = alpha values, y-axis = RMSE**

Task steps

- **Important Notes**

- Run each algorithm using the following hyper-parameter setup:
 - $\gamma = 0.99$
 - **Linear decay procedure** for ϵ (in ϵ -greedy action selection)
 - $\epsilon_{\text{start}} = 0.75$
 - $\epsilon_{\text{end}} = 0.001$
 - Decay linearly from episode 0 until $0.3 \times$ total number of episodes
 - After decay continue with ϵ_{end}